

CS310 Natural Language Processing

自然语言处理

Lecture 08 - Pretraining

Instructor: Yang Xu

主讲人：徐炀

xuyang@sustech.edu.cn

Overview


- Architecture
- Tokenization
- Pretraining
- Fine-Tuning
- Contextual Embedding and Word Sense

Overview

- **Architecture**
 - Autoregressive decoder: GPT
 - Autoencoding: BERT
 - Encoder-decoder: T5
 - General language modeling (GLM)
- Tokenization
- Pretraining
- Fine-Tuning
- Contextual Embedding and Word Sense

Background

- “pre-” means “before”
- Pretraining: train the model before applying it to a specific task
- Around 2017, state-of-the-art NLP = pretrained word embedding + LSTMs




Issues:

- The training data must be sufficient to cover broad aspects of language
- Most params are randomly initialized
- Word representations are not contextualized:
Ex. “movie” has fixed meaning no matter which sentence it appears in

figure from: <https://web.stanford.edu/class/archive/cs/cs224n/cs224n.1224/>

Pretraining the whole models



Ex. “movie” has dynamic meaning decided by its context

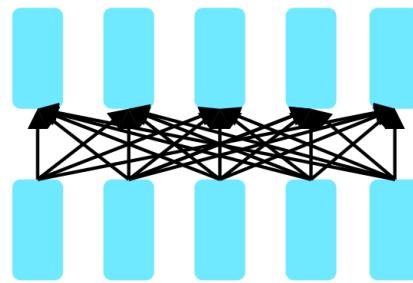
In modern NLP, especially since 2018, **all (or almost all) parameters** are initialized via pretraining

Pretraining:

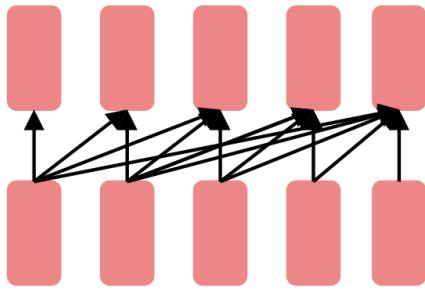
- Better contextualized representations of language
- “Warm up” model params with better initialization
- Make use of large text data

Architecture


- “Pretraining” can apply to all three types of neural architecture



Encoders



Decoders



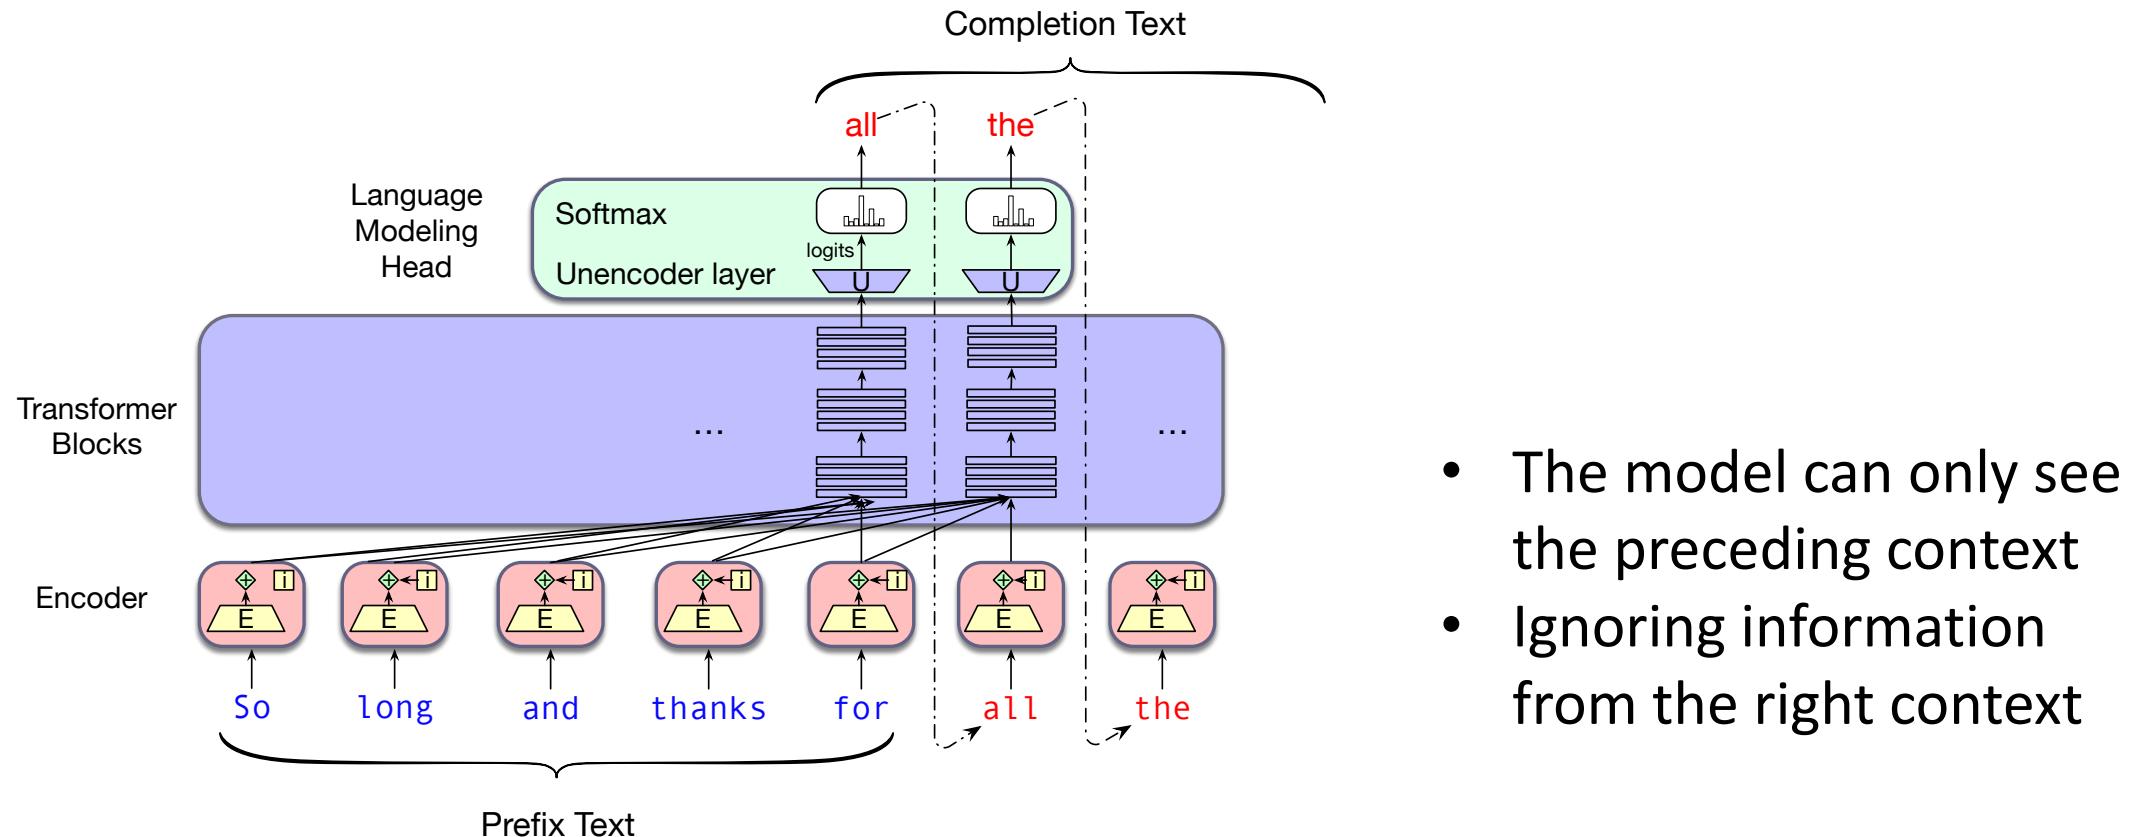
Encoder-decoders

Primary cases

自回归


Autoregressive

- Autoregressive = causal = left-to-right = decoder-only



Encoder Model

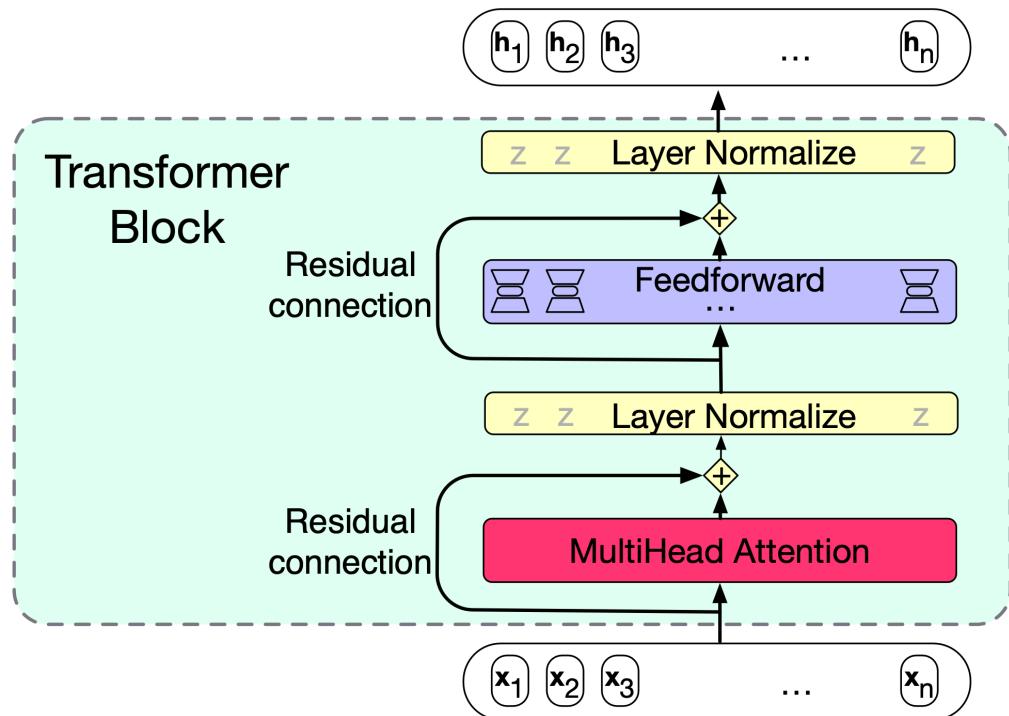
- Bidirectional Encoder Representations from Transformers (BERT) (Devlin et al., 2018)
- Its variants: **RoBERTa** (Liu et al., 2019), **ALBERT** (Lan et al., 2019) etc.
- Bidirectional encoders overcome this limitation by allowing self-attention to range over the entire input



a) A causal self-attention layer

b) A bidirectional self-attention layer

BERT Architecture



Uses the same self-attention as causal transformer

$$\begin{aligned} \mathbf{Q}_i &= \mathbf{X}\mathbf{W}_i^Q \\ \mathbf{K}_i &= \mathbf{X}\mathbf{W}_i^K \\ \mathbf{V}_i &= \mathbf{X}\mathbf{W}_i^V \end{aligned}$$

head_i
 $= \text{SelfAttention}(\mathbf{Q}_i, \mathbf{K}_i, \mathbf{V}_i)$

$$\begin{aligned} \mathbf{A} &= \text{MultiheadAttention}(\mathbf{X}) \\ &= (\mathbf{head}_1 \oplus \mathbf{head}_2 \dots \oplus \mathbf{head}_h) \mathbf{W}^O \end{aligned}$$

$$\text{SelfAttention}(Q, K, V) = \text{softmax} \left(\frac{QK^\top}{\sqrt{d_k}} \right) V$$

No future mask is used!

BERT: No Future Mask Used

q1·k1	q1·k2	q1·k3	q1·k4	q1·k5
q2·k1	q2·k2	q2·k3	q2·k4	q2·k5
q3·k1	q3·k2	q3·k3	q3·k4	q3·k5
q4·k1	q4·k2	q4·k3	q4·k4	q4·k5
q5·k1	q5·k2	q5·k3	q5·k4	q5·k5

N

The QK^T matrix contains **all pairs** of comparison between input queries and keys

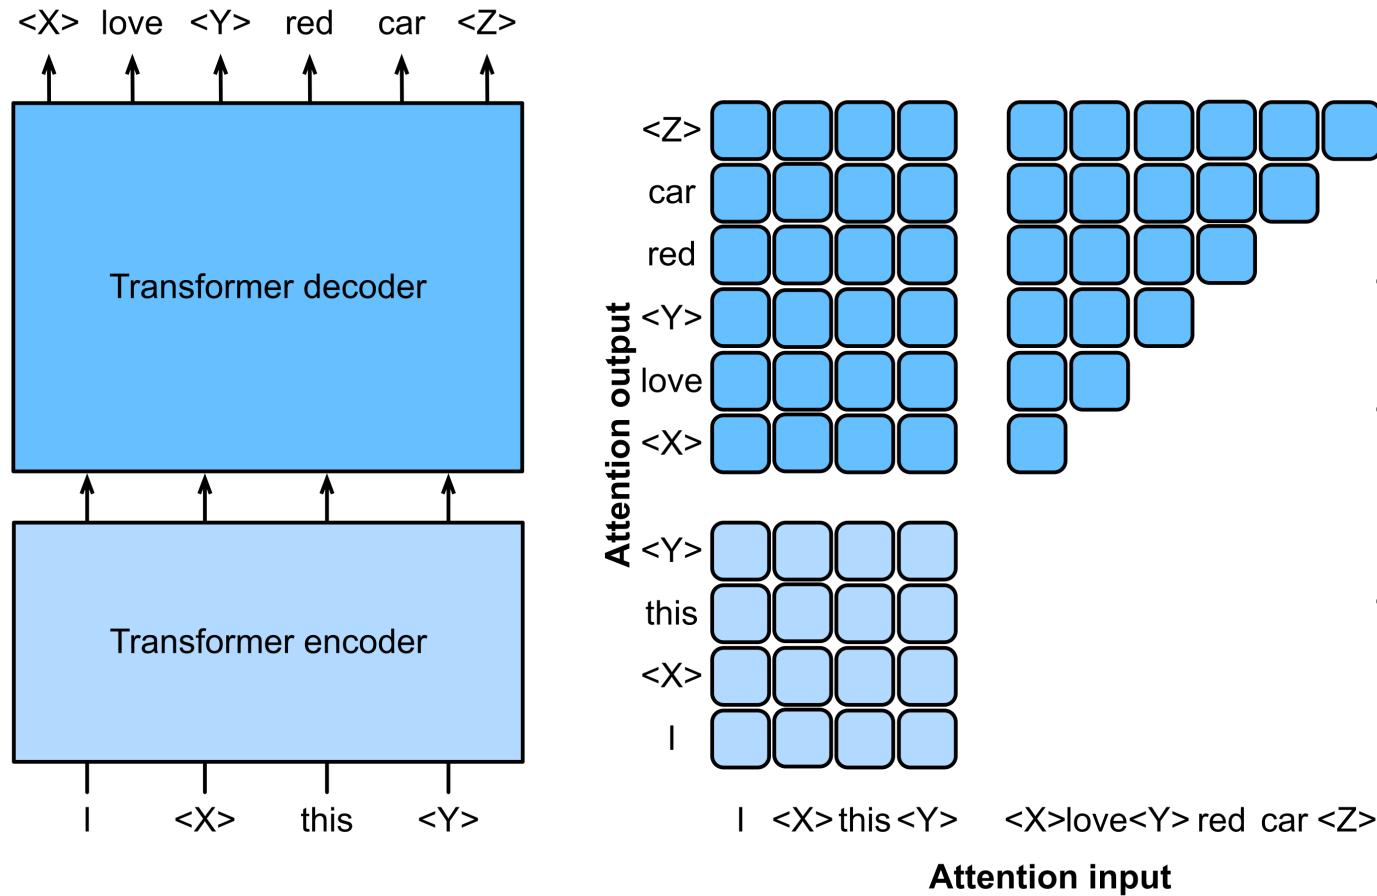
allowing the model to contextualize each token using *information from the entire input*

BERT: Hyper Parameters

- **BERT** (Devlin et al., 2018)
 - Vocabulary size: 30,000 tokens (from WordPiece algorithm)
 - Hidden layer size $d = 768$
 - 12 layers of transformer blocks, with 12 multihead attention layers each
 - Total params: about **100 M**
- **XLM-RoBERTa** (Liu et al., 2019)
 - Vocabulary size: 250,000 tokens (from SentencePiece Unigram LM algorithm)
 - Hidden layer size $d = 1024$; 24 layers of transformer blocks, with 16 multihead attention layers each
 - Total params: about **550 M**

Encoder-Decoder Model

- T5 (Google, 2019)

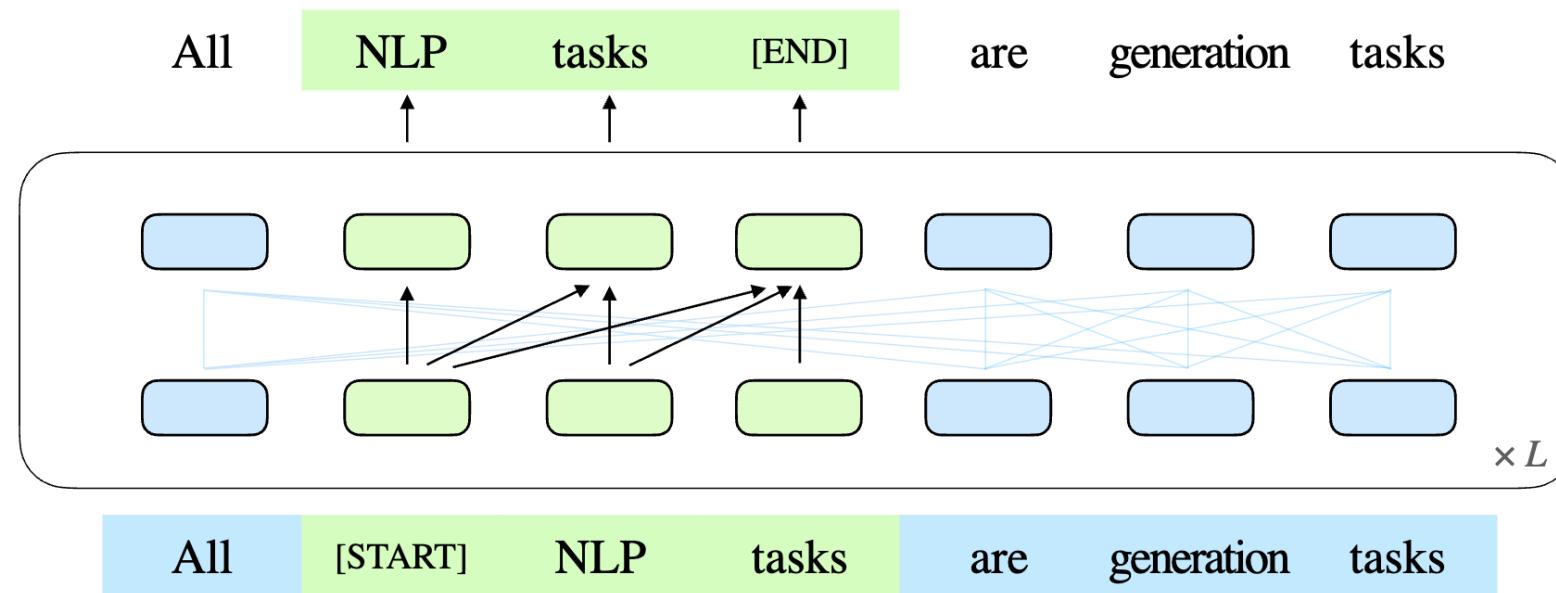


- Encoder self-attention (lower square), all input tokens attend to each other;
- Encoder–decoder cross-attention (upper rectangle), each target token attends to all input tokens
- Decoder self-attention (upper triangle), each target token attends to present and past target tokens only

General language modeling (GLM)

Du et al., 2022

- GLM: Randomly blank out spans of tokens and train the model to sequentially reconstruct the spans.



GLM

Du et al., 2022

 $x_1 \quad x_2 \quad x_3 \quad x_4 \quad x_5 \quad x_6$

(a) Sample spans from the input text

Part A: $x_1 \quad x_2 \quad [M] \quad x_4 \quad [M]$

Part B: $x_3 \quad x_5 \quad x_6$

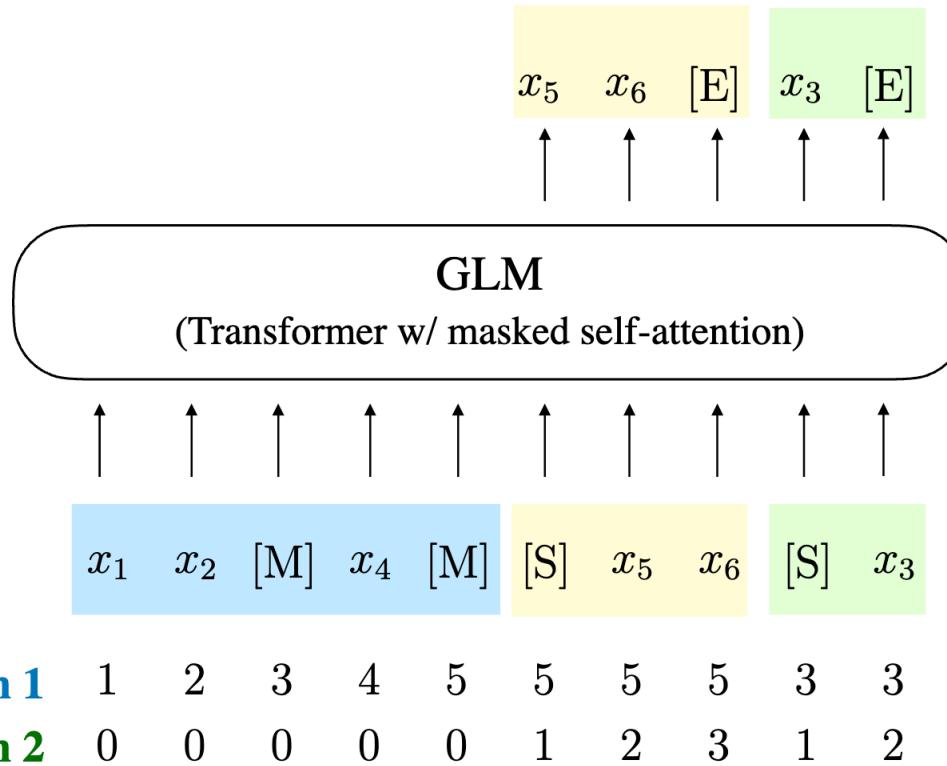
(b) Divide the input into Part A / Part B

The original text is $[x_1, x_2, x_3, x_4, x_5, x_6]$.
 Two spans $[x_3]$ and $[x_5, x_6]$ are sampled.

Replace the sampled spans with $[M]$ in Part A, and shuffle the spans in Part B

GLM

Du et al., 2022




(c) Generate the Part B spans autoregressively

Overview

- Architecture
- **Tokenization**
- Pretraining
- Fine-Tuning
- Contextual Embedding and Word Sense

General tokenization workflow



Here “model” means a **tokenizer model**

Example: **tday** => **td**, **##ay**

Figure source: <https://huggingface.co/learn/llm-course/en/chapter6/4?fw=pt>

Why subword is necessary

- Weakness of word-level tokenization:
- All *novel* words are mapped to UNK at testing time

	word	vocab mapping
Common words	hat	→ pizza (index)
	learn	→ tasty (index)
Variations	taaaaasty	→ UNK (index)
	laern	→ UNK (index)
misspellings		
novel items	Transformerify	→ UNK (index)

Subword-Level Tokenization

- Subword tokenization learns a vocabulary of subword tokens
- Split each word into a sequence of known subwords, for both training and testing

	word	→	vocab mapping
Common words	hat	→	hat
	learn	→	learn
Variations	taaaaasty	→	taa## aaa## sty
	laern	→	la## ern##
misspellings			
novel items	Transformerify	→	Transformer## ify

The hashtag “#” usage is from the WordPiece tokenization algorithm used in BERT

Byte-Pair Encoding (BPE) Tokenization

- A “bottom-up” algorithm: building the vocabulary by incrementally merging the most frequent pairs

Example corpus:

"hug", "pug", "pun", "bun", "hugs"

Initial vocabulary:

["b", "g", "h", "n", "p", "s", "u"]

with frequency:

("hug", 10), ("pug", 5), ("pun", 12), ("bun", 4), ("hugs", 5)

Example credits to: <https://huggingface.co/learn/llm-course/en/chapter6/5?fw=pt>

BPE Tokenization cont.

- Splitting each word into characters

```
("h" "u" "g", 10), ("p" "u" "g", 5), ("p" "u" "n", 12), ("b" "u" "n", 4), ("h" "u" "g" "s", 5)
```

Look at **pairs**, and **merge** the most frequent one: ("u", "g") -> "ug".

Vocabulary: ["b", "g", "h", "n", "p", "s", "u", "ug"]

Corpus: ("h" "ug", 10), ("p" "ug", 5), ("p" "u" "n", 12), ("b" "u" "n", 4), ("h" "ug" "s", 5)

BPE Tokenization cont.

- Given the new corpus, merge the next most frequent pair ("u", "n") -> "un"

Vocabulary: ["b", "g", "h", "n", "p", "s", "u", "ug", "un"]

Corpus: ("h" "ug", 10), ("p" "ug", 5), ("p" "un", 12), ("b" "un", 4), ("h" "ug" "s", 5)

- Next ("h", "ug") -> "hug"

Vocabulary: ["b", "g", "h", "n", "p", "s", "u", "ug", "un", "hug"]

Corpus: ("hug", 10), ("p" "ug", 5), ("p" "un", 12), ("b" "un", 4), ("hug" "s", 5)

- ...
- repeat until a desired vocabulary size is reached

Deal with Unicode characters

- GPT-2, RoBERTa, and recent Qwen, DeepSeek series use **Byte-level BPE**
- Process words not with Unicode characters, but with bytes
- Every character (Chinese, Japanese, Korean, emoji, ...) will be included, but not end up being unknown tokens

Overview

- Architecture
- Tokenization
- **Pretraining Task**
 - **Token level: Masked Language Models**
 - **Sequence level: Next Sentence Prediction**
 - **Training Method**
- Fine-Tuning
- Contextual Embedding and Word Sense

Pretraining of Bidirectional Encoders

- A new training scheme other than next-word-prediction:
- **Fill-in-Blank task -- close task**
- Predict a missing item given the rest of the sentence


Please turn _____ homework in

- instead of predicting the next word

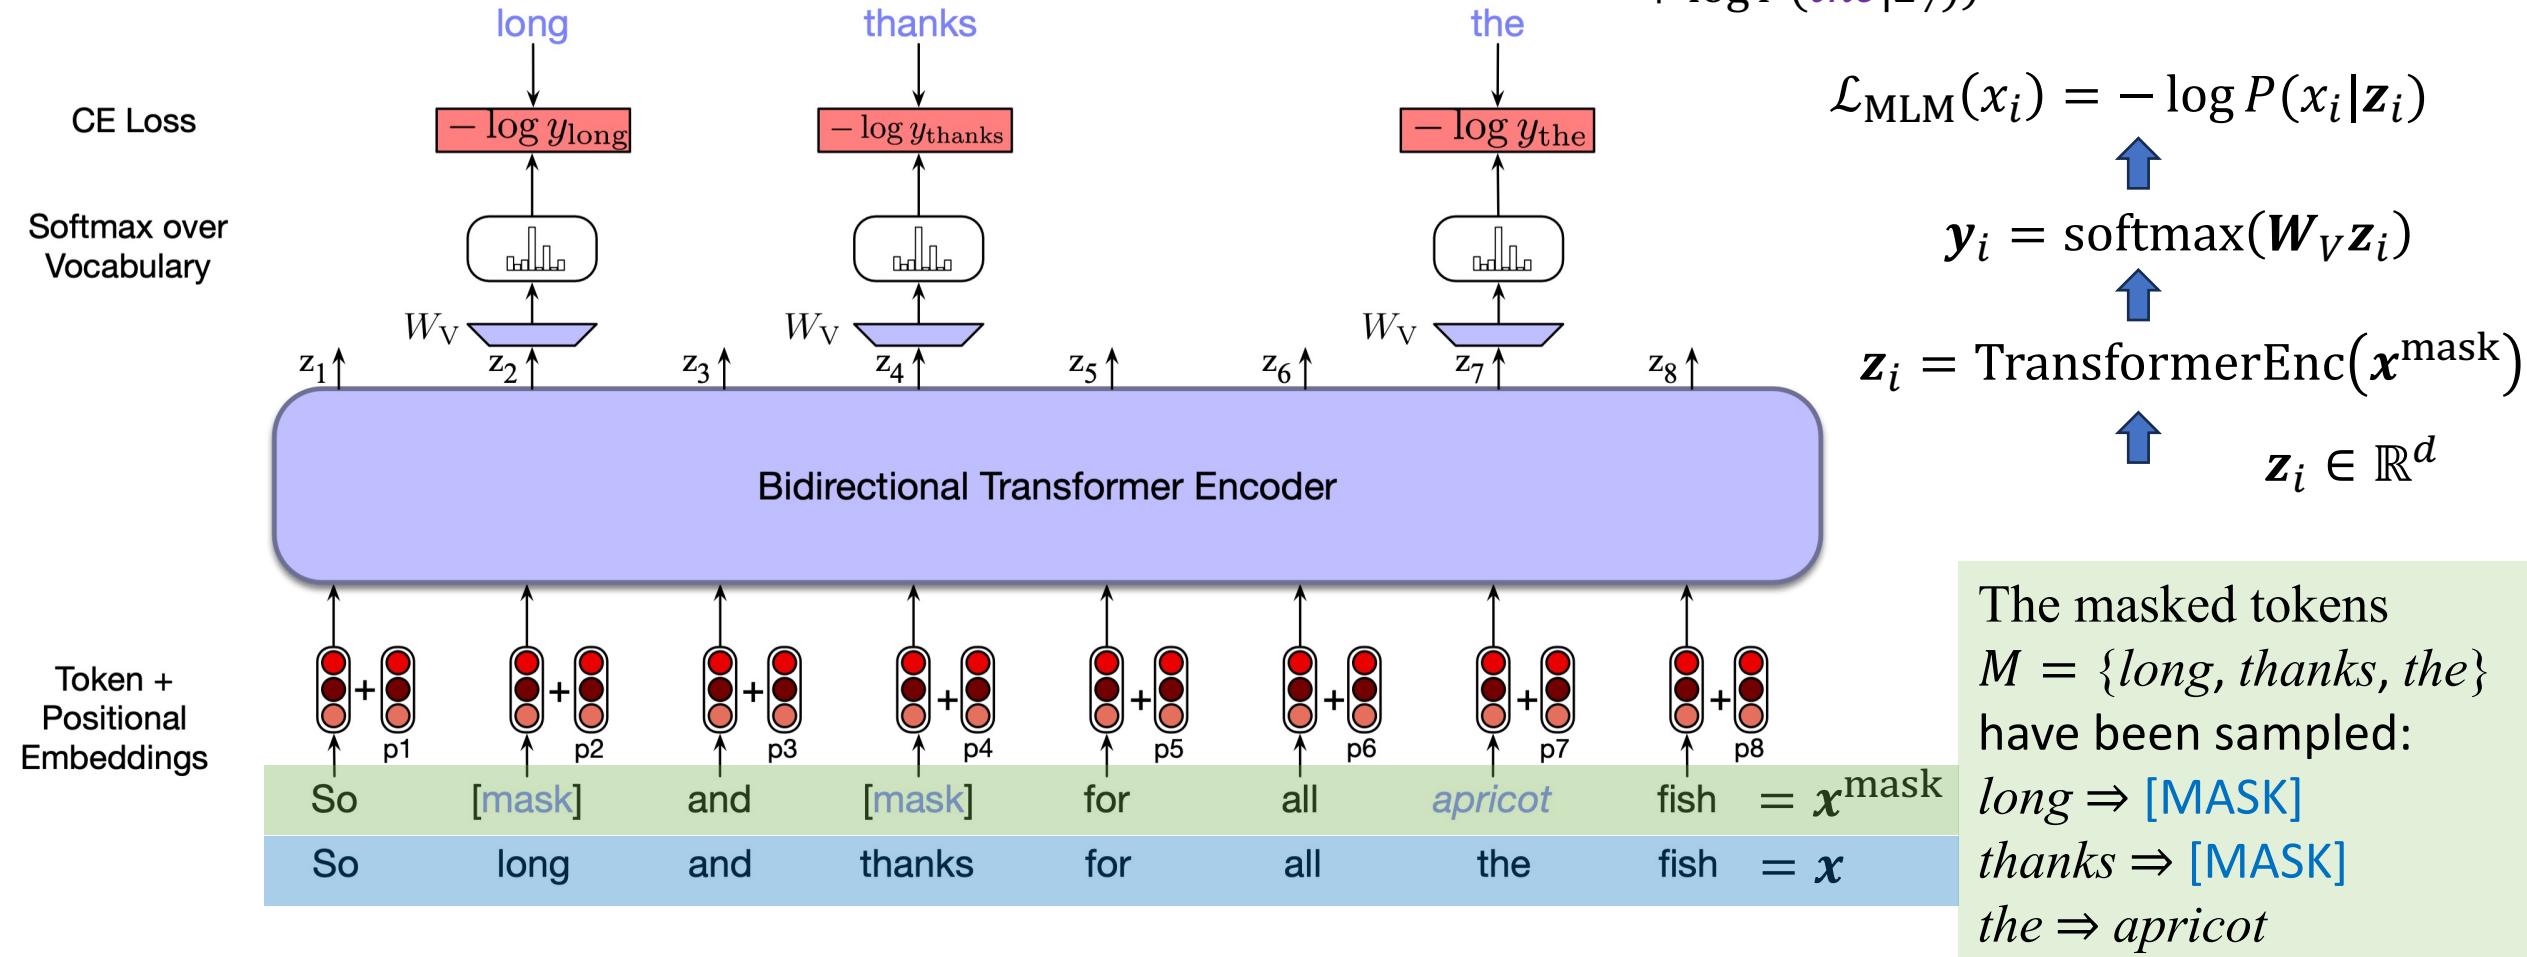
Please turn your homework _____

Masked Language Modeling

- The original approach used in BERT: Masked Language Modeling (MLM)
- **Randomly sample 15%** of the input tokens from each training sentence
- Among these sampled tokens:
 - \Rightarrow 80% are replaced with a special token **[MASK]**
 - \Rightarrow 10% are replaced with another randomly selected token
 - \Rightarrow 10% are left unchanged
- Why? Prevent the model from getting complacent and not building strong representations of non-masked words
- All tokens play a role in self-attention process, but only the sampled tokens are used for learning



MLM Overview



MLM Notes

- Only 15% of the input data are used for training
- Because only the sampled tokens in M play a role in the loss $\mathcal{L}_{MLM} = -\frac{1}{|M|} \sum_{i \in M} \log P(x_i | \mathbf{z}_i)$
- Other tokens play no role
- Thus, training BERT with MLM is inefficient
- Some members of BERT family use all examples for training, e.g., ELECTRA (Clark et al., 2020)

Next Sentence Prediction

- Some NLP applications involves determining the relationship between pairs of sentences, such as:
- **Paraphrase (改述)** detection: if two sentences have similar meanings
- **Entailment (蕴含)**: detect if s_1 entails s_2 , or contradict
- **Discourse coherence (话语连贯性)**: detect if two neighboring sentences form a coherent discourse

Paraphrase

- s1: "这个苹果很甜。"
(This apple is very sweet.)
- s2: "这个苹果味道很好。"
(This apple tastes really good.)

Entailment

- Premise: "The sky is blue."
- Hypothesis: "The sky is a shade of azure."

In this example, the hypothesis logically follows from the premise, indicating entailment.

Discourse coherence

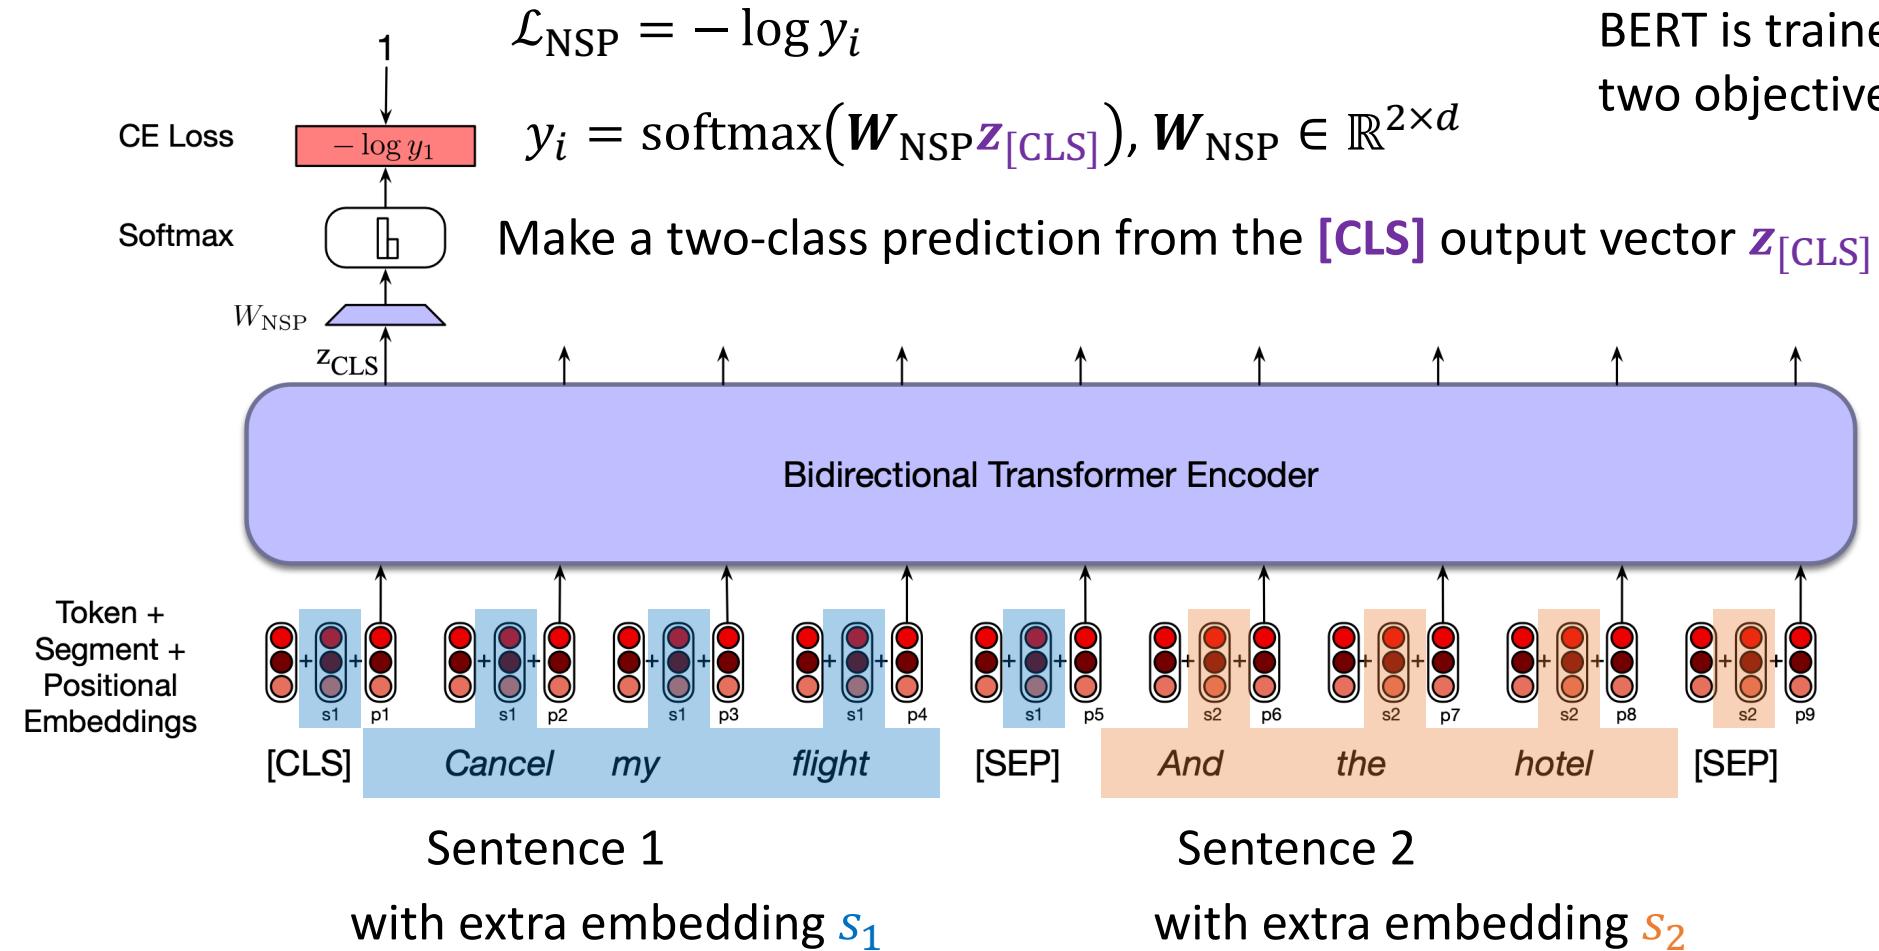
"我想要一杯咖啡。狗在睡觉。
" (I want a cup of coffee. The dog is sleeping.)

This is an negative example

Next Sentence Prediction

- Presented with pairs of sentences, the model's task is to predict whether each pair is an **ACTUAL** pair of adjacent sentences from the training corpus, or a pair of randomly sampled unrelated sentences
- BERT uses 50% of sentences in true pairs, and the other 50% random pairs
- Two new tokens to facilitate training:
- [CLS]**: prepended to input pair s_1, s_2 ("CLS" for "classification")
- [SEP]**: placed between s_1 and s_2 , and after s_2
- Finally, add extra embeddings to distinguish s_1 and s_2

Next Sentence Prediction



BERT is trained by combining the two objectives: $\mathcal{L}_{\text{MLM}} + \mathcal{L}_{\text{NSP}}$

Training Details of BERT

- Two models released (Devlin et al., 2018)
 - BERT-base: 12 layers (transformer blocks), 768-dim, 12 attention heads, 110 million params
 - BERT-large: 24 layers, 1024-dim, 16 attention heads, 340 million params
- Training data
 - BooksCorpus (800 million words)
 - English Wikipedia (2,500 million words)
- Trained on 64 TPUs for 4 days
 - Pretraining is impractical on a single GPU
- Yet, fine-tuning is practical and common on a single GPU
 - “Pretrain once, fine-tune many times”

Overview

- Architecture
- Tokenization
- Pretraining
- **Fine-Tuning**
- Contextual Embedding and Word Sense

Fine-Tuning

- **Fine-tuning:** add a small set of application-specific parameters on top of pretrained models
- Use *labeled data* to train these application-specific parameters
- Either *freeze* or make only *minimal* adjustments to the pretrained parameters
- Common applications:
 - Sequence classification
 - Pair-Wise sequence classification
 - Sequence labeling
 - Span-based operations (more advanced)

Sequence Classification

- An input sequence is represented with a single consolidated representation
- Recall RNN: use the final hidden state to stand for the entire sequence
- In transformer-based model: use an additional **vector** to represent the entire sequence ⇒ sometimes called **sentence embedding**
- In BERT, use the output vector of **[CLS]**, $\mathbf{z}_{[\text{CLS}]}$, plays this role
 - [CLS] is added to vocabulary and prepended to all input sequences during pretraining
- $\mathbf{z}_{[\text{CLS}]}$ serves as input to a **classifier head** -- a network classifier to make relevant predictions

Sequence Classification

Example: Sentiment classification

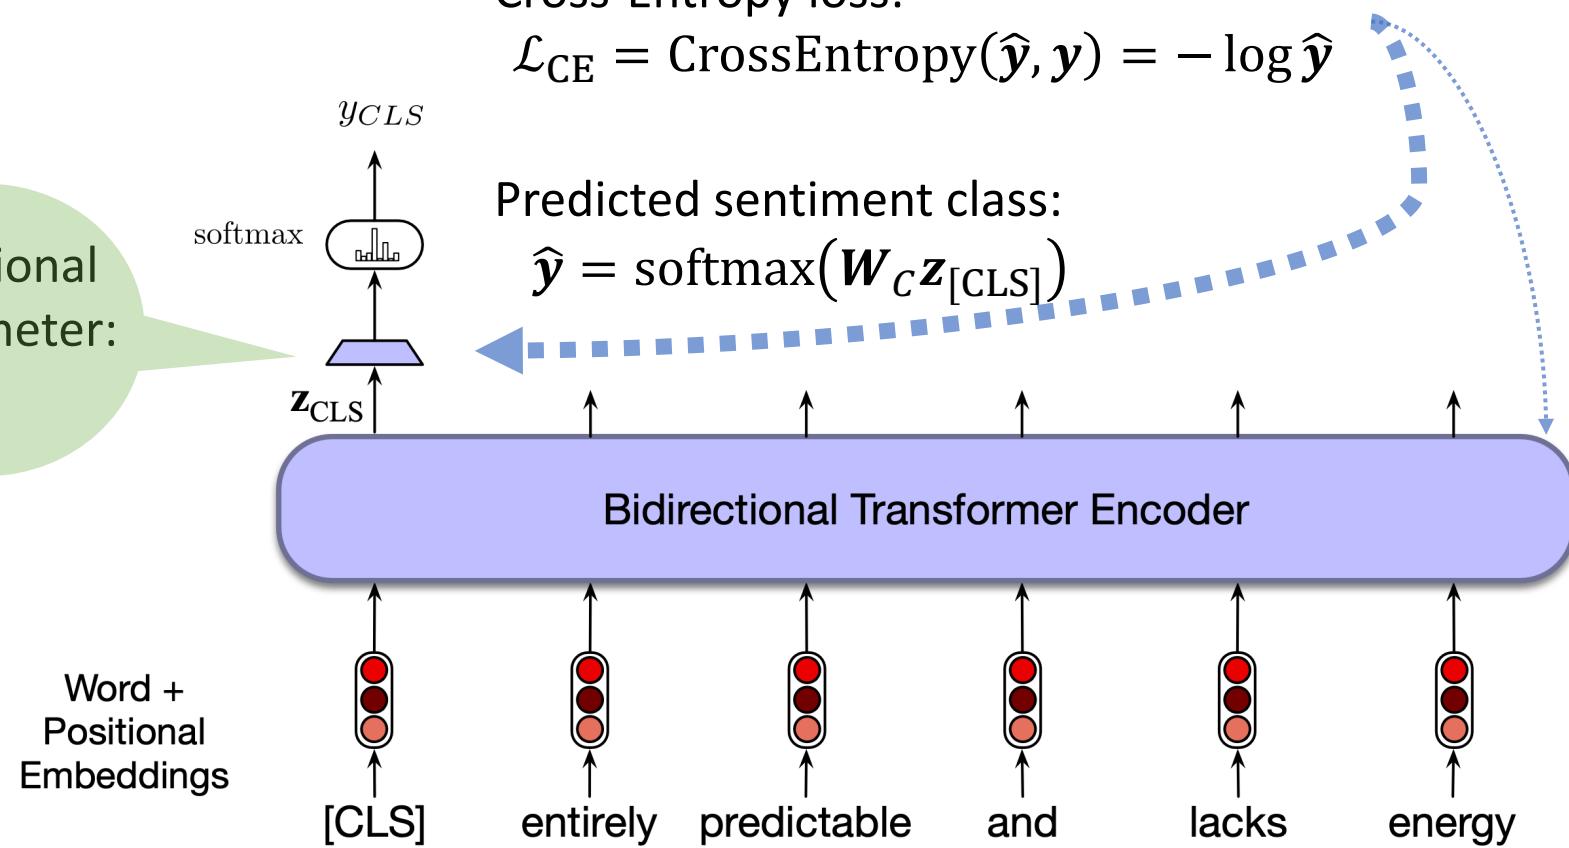
Additional parameter:
 W_C

Cross-Entropy loss:

$$\mathcal{L}_{CE} = \text{CrossEntropy}(\hat{\mathbf{y}}, \mathbf{y}) = -\log \hat{\mathbf{y}}$$

Predicted sentiment class:

$$\hat{\mathbf{y}} = \text{softmax}(\mathbf{W}_C \mathbf{z}_{[CLS]})$$




Difference from training a small neural classifier:

- Loss is mainly used to update param W_C
- Minimal changes are made to the encoder params: limited to updates over the **final few transformer layers**

Pair-Wise Sequence Classification

- **Paraphrase detection:** are sentence A and B paraphrases of each other?
- **Logical entailment:** does sentence A logically entail sentence B?
- **Discourse coherence:** how coherent is sentence B as a follow-on to sentence A?
- Fine-tuning these tasks proceeds just as with the next sentence prediction (NSP) pretraining task



Example: entailment classification

- Dataset: Multi-Genre Natural Language Inference (MultiNLI)
- Pairs of sentences mapped to one of 3 labels: {entail, contradicts, neutral}

Entails

- a: I'm confused.
b: Not all of it is very clear to me.

Neutral

- a: Jon walked back to the town to the smithy.
b: Jon traveled back to his hometown.

Contradicts


- a: Tourist Information offices can be very helpful.
b: Tourist Information offices are never of any help.

$$y_i = \text{softmax}(\mathbf{W}\mathbf{z}_{[\text{CLS}]}) \in \mathbb{R}^3$$

(Williams et al., 2018)

Sequence Labeling

- The final output vector z_i corresponding to **each input token** x_i is passed to a classifier to produce a probability distribution over possible labels
- Example: POS tagging, BIO-based NER



Additional param: $W_K \in \mathbb{R}^{k \times d}$
for k possible tags

with greedy approach:
 $\hat{y}_i = \text{softmax}(W_K z_i)$
 $y_i = \arg \max_k \hat{y}_i$


More advanced: pass \hat{y}_i to a conditional random field (CRF) layer and use viterbi decoding

Overview

- Architecture
- Tokenization
- Pretraining
- Fine-Tuning
- **Contextual Embedding and Word Sense**

Contextual Embeddings

- The output of a BERT-style model is a contextual embedding vector z_i for each input token x_i
 - z_i represents some aspects of the meanings of x_i
 - Sometimes, instead of just using the final layer, we can average the z_i from the last four layers



Contextual Embeddings

- Static embedding (e.g., word2vec) ⇒ meaning of a word *types*
 - a type is a static entry in the vocabulary
- Contextual embedding ⇒ meaning of word *instances*
 - Instances of a particular word type in a particular context
- Thus, contextual embeddings are useful in linguistic tasks that require models of word meaning
 - “I would like some **orange** juice” ⇒ fruit
 - “Paint this part **orange**” ⇒ color

Word Sense

- Words are ambiguous: same word can be used to mean different things
- **Polysemous** (多义词), Geek “many senses”
- A **word sense** is a discrete representation of *one meaning* of a word

mouse¹ : a *mouse* controlling a computer system in 1968.

mouse² : a quiet animal like a *mouse*

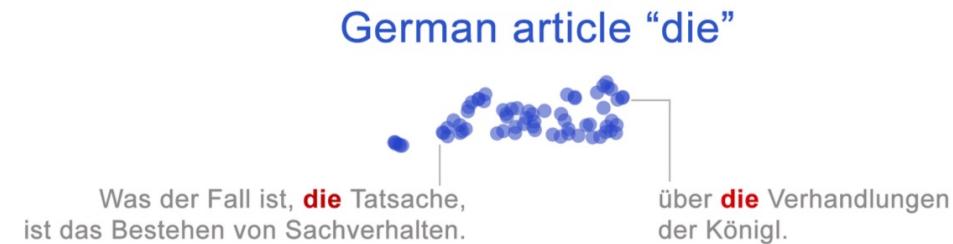
bank¹ : ...a *bank* can hold the investments in a custodial account ...

bank² : ...as agriculture burgeons on the east *bank*, the river ...

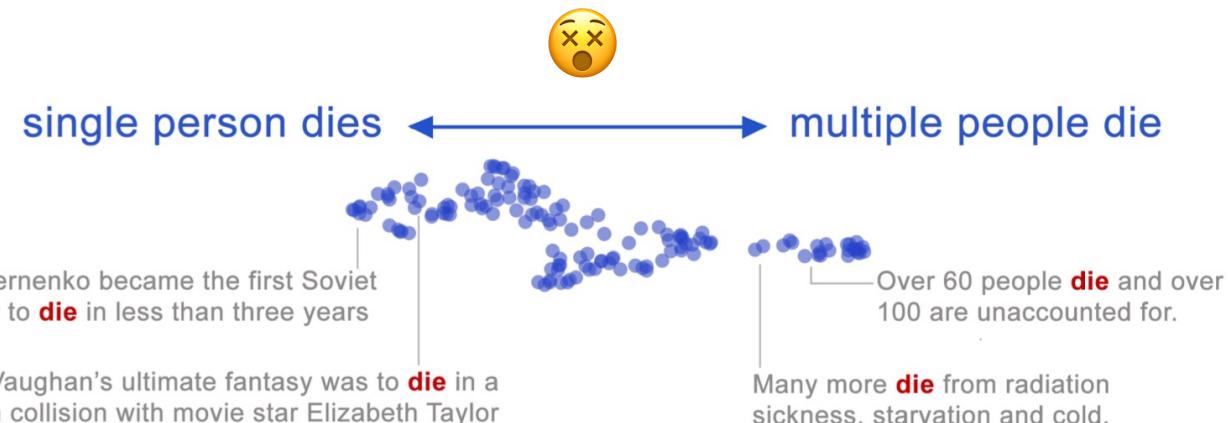
- The senses can be visualized geometrically by contextual embeddings

Visualize Word Senses

Figure from Coenen et al. (2019)

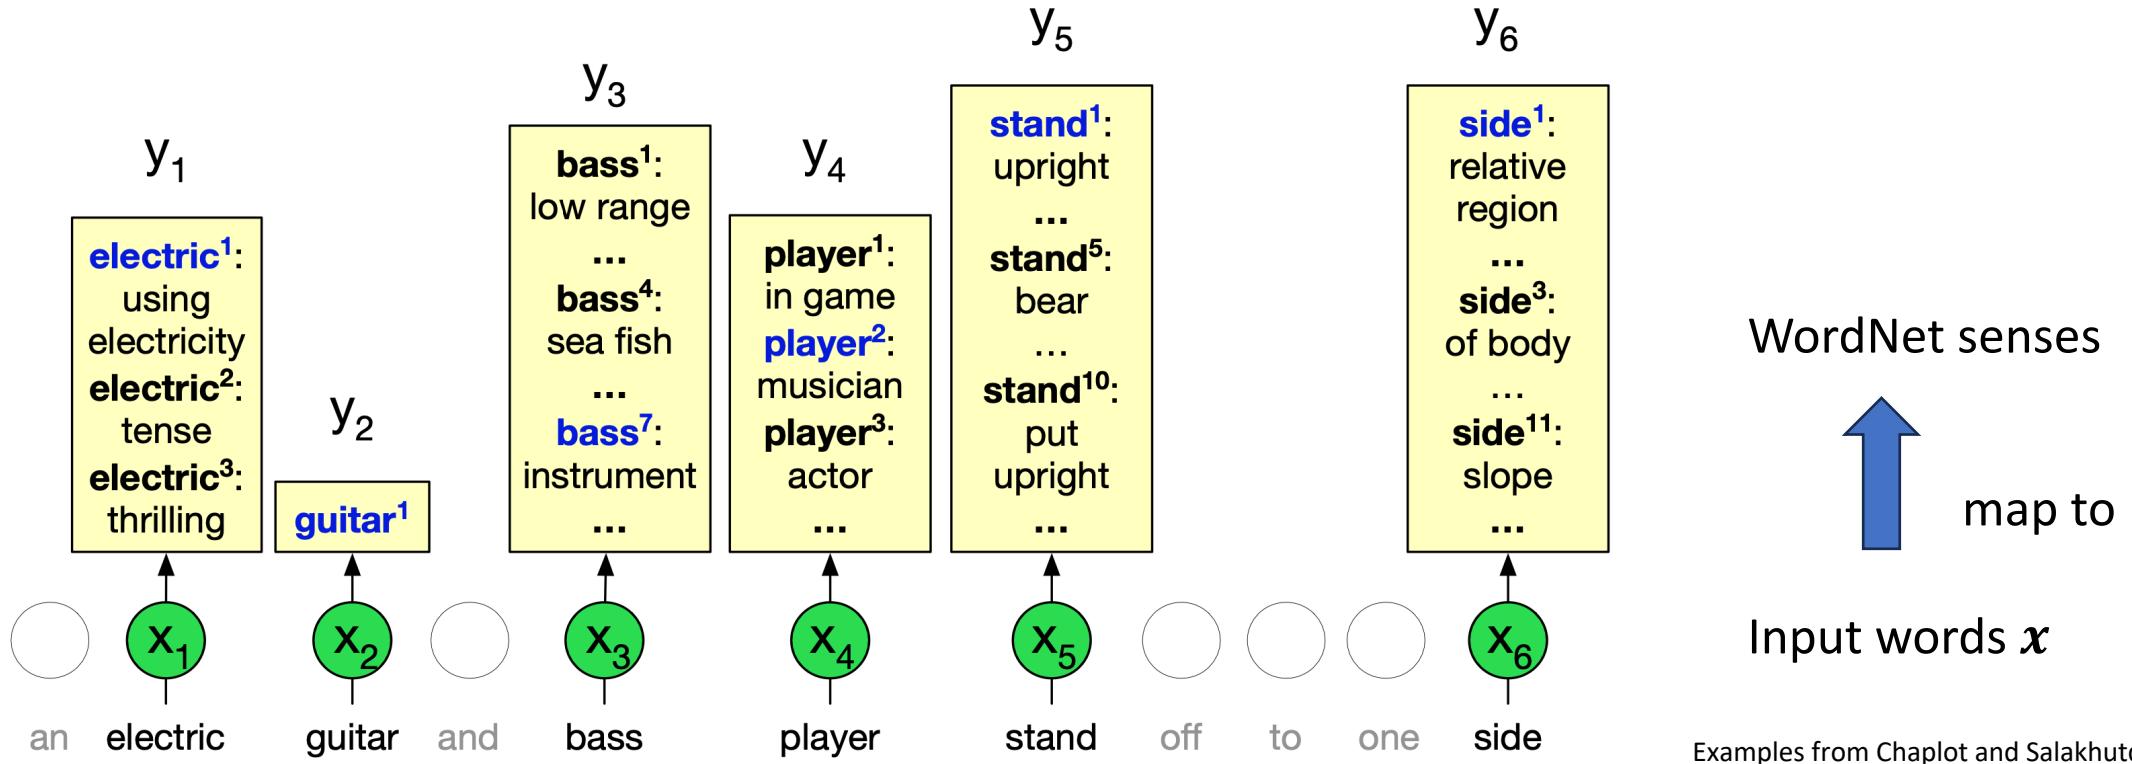


- Dictionary and thesauruses like WordNet give **discrete** lists of senses
- Embeddings (whether contextual or static) provides a **continuous high-dimensional** view of meaning



Word Sense Disambiguation

- **Word Sense Disambiguation:** the task of selecting the correct sense for a word
- Takes as input a word in context and a fixed inventory of potential word senses (like the ones in WordNet) and outputs the correct sense in context



WSD Algorithm: 1-nearest-neighbor

- At training time: pass a sense-labeled dataset through any contextual embedding (e.g., BERT) \Rightarrow vector v_i for each token i
- For each sense s of any word, and for each of the n tokens of that sense, produce a **sense embedding** v_s by averaging the n contextual embeddings:

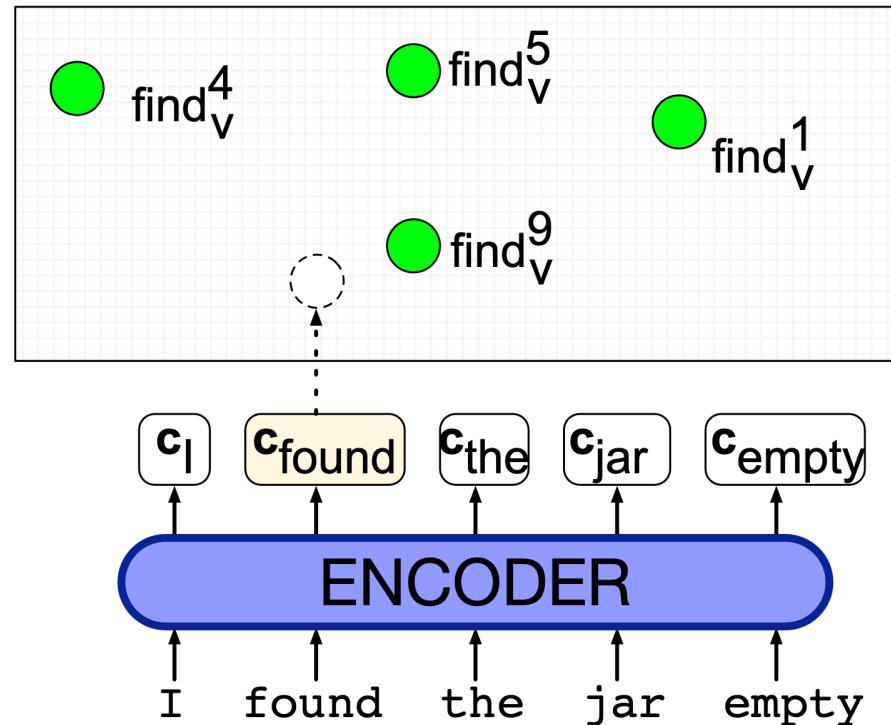
$$v_s = \frac{1}{n} \sum_i v_i \quad \forall v_i \in \text{tokens}(s)$$

- At test time, given a token of target word t , compute its contextual embedding t and choose its nearest neighbor sense from the training set:

$$\text{sense}(t) = \arg \max_{s \in \text{senses}(t)} \cos(t, v_s)$$

WSD 1-nearest-neighbor Example

- Contextual embedding for $found$ c_{found} is computed, and the nearest neighbor sense \mathbf{find}_v^9 is chosen



Examples from Loureiro and Jorge, 2019

Word Similarity is Tricky

- **Fact:** Contextual embeddings for all words are *extremely similar*
- **Fact:** BERT embeddings of any two randomly chosen words will have extremely high cosines $\approx 1 \Rightarrow$ All word vectors tend to point in the same direction
- A property known as **anisotropy** (各向异性)
- **Anisotropy** \triangleq the expected cosine similarity of any pair of words in a corpus (Ethayarajh, 2019)
- If all vectors are uniformly distributed, then the expected cosine should be 0, which we call **isotropy** (各向同性)
- Cause of anisotropy: cosine measures are dominated by a small number of rogue dimensions that have very large magnitudes and high variance (Timkey and van Schijndel, 2021)

Solution to Anisotropy

- Standardizing (z-scoring) the vectors, i.e., subtracting the mean and dividing by variance

$$\mu = \frac{1}{|C|} \sum_{x \in C} x \quad \sigma = \sqrt{\frac{1}{|C|} \sum_{x \in C} (x - \mu)^2} \quad z = \frac{x - \mu}{\sigma}$$

- Remaining problem: cosine tends to underestimate similarity of word meanings for very frequent words (according to human judgements) (Zhou et al., 2022)

Reference

- Devlin, J., Chang, M. W., Lee, K., & Toutanova, K. (2018). Bert: Pre-training of deep bidirectional transformers for language understanding. *arXiv preprint arXiv:1810.04805*.
- Liu, Y., Ott, M., Goyal, N., Du, J., Joshi, M., Chen, D., ... & Stoyanov, V. (2019). Roberta: A robustly optimized bert pretraining approach. *arXiv preprint arXiv:1907.11692*.
- Lan, Z., Chen, M., Goodman, S., Gimpel, K., Sharma, P., & Soricut, R. (2019). Albert: A lite BERT for self-supervised learning of language representations. *arXiv preprint arXiv:1909.11942*.
- Ethayarajh, K. (2019). How contextual are contextualized word representations? Comparing the geometry of BERT, ELMo, and GPT-2 embeddings. *arXiv preprint arXiv:1909.00512*.
- Timkey, W., & Van Schijndel, M. (2021). All bark and no bite: Rogue dimensions in transformer language models obscure representational quality. *arXiv preprint arXiv:2109.04404*.
- Zhou, K., Ethayarajh, K., Card, D., & Jurafsky, D. (2022). Problems with cosine as a measure of embedding similarity for high frequency words. *arXiv preprint arXiv:2205.05092*.